

Keras -- MLPs on MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [0]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic_plot(x, vy, ty, ax, colors=['b']):
    fig, ax = plt.subplots(1,1)
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [8]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

In [10]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.]
```

[illegible]

[illegible]

In [11]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/
```

```

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [0]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [0]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer

```

```
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 49us/step - loss: 1.2935 - acc: 0.6829 -
val_loss: 0.8171 - val_acc: 0.8312
Epoch 2/20
60000/60000 [=====] - 2s 35us/step - loss: 0.7213 - acc: 0.8385 -
val_loss: 0.6099 - val_acc: 0.8623
Epoch 3/20
60000/60000 [=====] - 2s 35us/step - loss: 0.5904 - acc: 0.8584 -
val_loss: 0.5275 - val_acc: 0.8741
Epoch 4/20
60000/60000 [=====] - 2s 35us/step - loss: 0.5278 - acc: 0.8682 -
val_loss: 0.4815 - val_acc: 0.8814
Epoch 5/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4897 - acc: 0.8747 -
val_loss: 0.4515 - val_acc: 0.8870
Epoch 6/20
34432/60000 [=====>.....] - ETA: 0s - loss: 0.4697 - acc: 0.877360000/60000 [=====] - 2s 35us/step - loss: 0.4635 - acc: 0.8799 - val_loss: 0.4303 - val_acc: 0.8898
Epoch 7/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4442 - acc: 0.8837 -
val_loss: 0.4138 - val_acc: 0.8917
Epoch 8/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4290 - acc: 0.8866 -
val_loss: 0.4010 - val_acc: 0.8952
Epoch 9/20
```

```

Epoch 9/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4169 - acc: 0.8894 -
val_loss: 0.3909 - val_acc: 0.8971
Epoch 10/20
60000/60000 [=====] - 2s 35us/step - loss: 0.4067 - acc: 0.8911 -
val_loss: 0.3818 - val_acc: 0.8994
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3982 - acc: 0.8926 -
val_loss: 0.3743 - val_acc: 0.9006
Epoch 12/20
1792/60000 [.....] - ETA: 1s - loss: 0.4077 - acc: 0.889560000/60000 [==
=====] - 2s 35us/step - loss: 0.3908 - acc: 0.8948 - val_loss: 0.3681 - val
_acc: 0.9020
Epoch 13/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3844 - acc: 0.8960 -
val_loss: 0.3624 - val_acc: 0.9039
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3786 - acc: 0.8972 -
val_loss: 0.3575 - val_acc: 0.9046
Epoch 15/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3735 - acc: 0.8981 -
val_loss: 0.3528 - val_acc: 0.9058
Epoch 16/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3689 - acc: 0.8993 -
val_loss: 0.3490 - val_acc: 0.9062
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3648 - acc: 0.9004 -
val_loss: 0.3455 - val_acc: 0.9066
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3610 - acc: 0.9016 -
val_loss: 0.3419 - val_acc: 0.9077
Epoch 19/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3575 - acc: 0.9024 -
val_loss: 0.3389 - val_acc: 0.9088
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.3544 - acc: 0.9032 -
val_loss: 0.3362 - val_acc: 0.9093

```

In [0]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3362289469957352
Test accuracy: 0.9093

MLP + Sigmoid activation + SGDOptimizer

In [0]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 42us/step - loss: 2.2708 - acc: 0.2169 -
val_loss: 2.2197 - val_acc: 0.2768
Epoch 2/20
60000/60000 [=====] - 2s 42us/step - loss: 2.1739 - acc: 0.4237 -
val_loss: 2.1143 - val_acc: 0.4877
Epoch 3/20
60000/60000 [=====] - 2s 42us/step - loss: 2.0506 - acc: 0.5485 -
val_loss: 1.9659 - val_acc: 0.5524
Epoch 4/20
60000/60000 [=====] - 3s 42us/step - loss: 1.8796 - acc: 0.6135 -
val_loss: 1.7673 - val_acc: 0.6481
Epoch 5/20
60000/60000 [=====] - 3s 42us/step - loss: 1.6674 - acc: 0.6659 -
val_loss: 1.5414 - val_acc: 0.7205
Epoch 6/20
5376/60000 [=>.....] - ETA: 2s - loss: 1.5430 - acc: 0.698160000/60000 [==
=====] - 2s 41us/step - loss: 1.4449 - acc: 0.7125 - val_loss: 1.3233 - val
_acc: 0.7513
Epoch 7/20
60000/60000 [=====] - 2s 41us/step - loss: 1.2442 - acc: 0.7466 -
val_loss: 1.1406 - val_acc: 0.7599
Epoch 8/20
60000/60000 [=====] - 2s 41us/step - loss: 1.0815 - acc: 0.7722 -
val_loss: 0.9974 - val_acc: 0.7968
Epoch 9/20
60000/60000 [=====] - 2s 41us/step - loss: 0.9544 - acc: 0.7940 -
val_loss: 0.8867 - val_acc: 0.8060
Epoch 10/20
60000/60000 [=====] - 2s 41us/step - loss: 0.8556 - acc: 0.8082 -
val_loss: 0.7984 - val_acc: 0.8227
Epoch 11/20
31232/60000 [=====>.....] - ETA: 1s - loss: 0.7920 - acc: 0.820760000/60000 [==
=====] - 2s 42us/step - loss: 0.7776 - acc: 0.8225 - val_loss: 0.7292 - val
_acc: 0.8335
Epoch 12/20
60000/60000 [=====] - 2s 41us/step - loss: 0.7154 - acc: 0.8333 -
val_loss: 0.6741 - val_acc: 0.8422
Epoch 13/20
60000/60000 [=====] - 2s 41us/step - loss: 0.6650 - acc: 0.8412 -
val_loss: 0.6282 - val_acc: 0.8500
Epoch 14/20
```



```

60000/60000 [=====] - 2s 41us/step - loss: 0.6237 - acc: 0.8485 -
val_loss: 0.5906 - val_acc: 0.8585
Epoch 15/20
60000/60000 [=====] - 2s 41us/step - loss: 0.5893 - acc: 0.8546 -
val_loss: 0.5591 - val_acc: 0.8616
Epoch 16/20
34432/60000 [=====>.....] - ETA: 0s - loss: 0.5691 - acc: 0.857860000/60000 [==
=====] - 2s 41us/step - loss: 0.5606 - acc: 0.8599 - val_loss: 0.5329 - val
_acc: 0.8672
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.5361 - acc: 0.8641 -
val_loss: 0.5095 - val_acc: 0.8703
Epoch 18/20
60000/60000 [=====] - 2s 42us/step - loss: 0.5151 - acc: 0.8676 -
val_loss: 0.4904 - val_acc: 0.8736
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.4969 - acc: 0.8710 -
val_loss: 0.4732 - val_acc: 0.8782
Epoch 20/20
60000/60000 [=====] - 2s 42us/step - loss: 0.4810 - acc: 0.8739 -
val_loss: 0.4583 - val_acc: 0.8816

```

In [0]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.4582893396139145
Test accuracy: 0.8816

In [0]:

```

w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

MLP + Sigmoid activation + ADAM

In [0]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 55us/step - loss: 0.5308 - acc: 0.8636 -
val_loss: 0.2550 - val_acc: 0.9259
Epoch 2/20
53120/60000 [=====>....] - ETA: 0s - loss: 0.2244 - acc: 0.934060000/60000 [==
=====] - 3s 51us/step - loss: 0.2205 - acc: 0.9351 - val_loss: 0.1946 - val
_acc: 0.9417
Epoch 3/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1650 - acc: 0.9512 -
val_loss: 0.1421 - val_acc: 0.9570
Epoch 4/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1279 - acc: 0.9614 -
val_loss: 0.1238 - val_acc: 0.9645
Epoch 5/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1010 - acc: 0.9704 -
val_loss: 0.1029 - val_acc: 0.9693
Epoch 6/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0801 - acc: 0.9763 -
val_loss: 0.0877 - val_acc: 0.9725
Epoch 7/20
4480/60000 [=>.....] - ETA: 2s - loss: 0.0632 - acc: 0.979560000/60000 [==
=====] - 3s 51us/step - loss: 0.0645 - acc: 0.9809 - val_loss: 0.0831 - val
_acc: 0.9751
Epoch 8/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0519 - acc: 0.9842 -
val_loss: 0.0724 - val_acc: 0.9780
Epoch 9/20
```

```

60000/60000 [=====] - 3s 51us/step - loss: 0.0433 - acc: 0.9872 -
val_loss: 0.0714 - val_acc: 0.9786
Epoch 10/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0347 - acc: 0.9898 -
val_loss: 0.0695 - val_acc: 0.9776
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0268 - acc: 0.9930 -
val_loss: 0.0659 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0219 - acc: 0.9944 -
val_loss: 0.0642 - val_acc: 0.9809
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0180 - acc: 0.9953 -
val_loss: 0.0677 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0133 - acc: 0.9970 -
val_loss: 0.0647 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0114 - acc: 0.9975 -
val_loss: 0.0628 - val_acc: 0.9812
Epoch 16/20
57344/60000 [=====>..] - ETA: 0s - loss: 0.0085 - acc: 0.998260000/60000 [==
=====] - 3s 50us/step - loss: 0.0085 - acc: 0.9982 - val_loss: 0.0666 - val
_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0070 - acc: 0.9986 -
val_loss: 0.0643 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0061 - acc: 0.9986 -
val_loss: 0.0656 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0055 - acc: 0.9988 -
val_loss: 0.0811 - val_acc: 0.9774
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0038 - acc: 0.9992 -
val_loss: 0.0723 - val_acc: 0.9818

```

In [0]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06385514608082886

Test accuracy: 0.9824





In [0]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)
```

MLP + ReLU +SGD

In [0]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.062)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.125)$ 
# out =>  $\sigma=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 
```

```

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```

model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 67us/step - loss: 0.7579 - acc: 0.7812 -
val_loss: 0.3951 - val_acc: 0.8921
Epoch 2/20
60000/60000 [=====] - 4s 64us/step - loss: 0.3535 - acc: 0.8998 -
val_loss: 0.3040 - val_acc: 0.9153
Epoch 3/20
60000/60000 [=====] - 4s 64us/step - loss: 0.2900 - acc: 0.9172 -
val_loss: 0.2648 - val_acc: 0.9253
Epoch 4/20
60000/60000 [=====] - 4s 60us/step - loss: 0.2558 - acc: 0.9269 -
val_loss: 0.2393 - val_acc: 0.9316
Epoch 5/20
60000/60000 [=====] - 4s 58us/step - loss: 0.2324 - acc: 0.9340 -
val_loss: 0.2210 - val_acc: 0.9371
Epoch 6/20
60000/60000 [=====] - 4s 64us/step - loss: 0.2144 - acc: 0.9391 -
val_loss: 0.2072 - val_acc: 0.9400
Epoch 7/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1995 - acc: 0.9443 -
val_loss: 0.1957 - val_acc: 0.9444
Epoch 8/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1872 - acc: 0.9476 -
val_loss: 0.1848 - val_acc: 0.9456
Epoch 9/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1763 - acc: 0.9507 -
val_loss: 0.1771 - val_acc: 0.9488
Epoch 10/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1668 - acc: 0.9539 -
val_loss: 0.1682 - val_acc: 0.9506
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1585 - acc: 0.9560 -
val_loss: 0.1623 - val_acc: 0.9518
Epoch 12/20
60000/60000 [=====] - 7s 113us/step - loss: 0.1511 - acc: 0.9577 -
val_loss: 0.1560 - val_acc: 0.9543
Epoch 13/20
60000/60000 [=====] - 7s 115us/step - loss: 0.1443 - acc: 0.9596 -
val_loss: 0.1517 - val_acc: 0.9557
Epoch 14/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1379 - acc: 0.9615 -
val_loss: 0.1474 - val_acc: 0.9572
Epoch 15/20

```

```

60000/60000 [=====] - 4s 66us/step - loss: 0.1323 - acc: 0.9628 -
val_loss: 0.1429 - val_acc: 0.9580
Epoch 16/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1270 - acc: 0.9645 -
val_loss: 0.1371 - val_acc: 0.9598
Epoch 17/20
60000/60000 [=====] - 7s 110us/step - loss: 0.1221 - acc: 0.9661 -
val_loss: 0.1351 - val_acc: 0.9602
Epoch 18/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1177 - acc: 0.9671 -
val_loss: 0.1309 - val_acc: 0.9618
Epoch 19/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1136 - acc: 0.9685 -
val_loss: 0.1263 - val_acc: 0.9631
Epoch 20/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1094 - acc: 0.9694 -
val_loss: 0.1241 - val_acc: 0.9631

```

In [0]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12405014228336513

Test accuracy: 0.9631





In [0]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [0]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 121us/step - loss: 0.2341 - acc: 0.9295 -
val_loss: 0.1165 - val_acc: 0.9652
Epoch 2/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0878 - acc: 0.9729 -
val_loss: 0.0883 - val_acc: 0.9720
Epoch 3/20
60000/60000 [=====] - 5s 75us/step - loss: 0.0544 - acc: 0.9825 -
val_loss: 0.0860 - val_acc: 0.9729
Epoch 4/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0354 - acc: 0.9885 -
val_loss: 0.0699 - val_acc: 0.9797
Epoch 5/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0266 - acc: 0.9914 -
val_loss: 0.0720 - val_acc: 0.9788
Epoch 6/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0200 - acc: 0.9941 -
val_loss: 0.0696 - val_acc: 0.9803
Epoch 7/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0155 - acc: 0.9951 -
val_loss: 0.0640 - val_acc: 0.9829
Epoch 8/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0140 - acc: 0.9952 -
val_loss: 0.0848 - val_acc: 0.9792
Epoch 9/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0143 - acc: 0.9952 -
val_loss: 0.0837 - val_acc: 0.9796
Epoch 10/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0128 - acc: 0.9958 -
val_loss: 0.0946 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0081 - acc: 0.9974 -
val_loss: 0.0682 - val_acc: 0.9826
Epoch 12/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0121 - acc: 0.9959 -
val_loss: 0.0793 - val_acc: 0.9816
Epoch 13/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0107 - acc: 0.9963 -
val_loss: 0.0746 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0113 - acc: 0.9960 -
val_loss: 0.0813 - val_acc: 0.9816
Epoch 15/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0058 - acc: 0.9982 -
val_loss: 0.0770 - val_acc: 0.9842
Epoch 16/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0040 - acc: 0.9987 -
val_loss: 0.0930 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0119 - acc: 0.9959 -
val_loss: 0.0813 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0105 - acc: 0.9966 -
val_loss: 0.1000 - val_acc: 0.9803
```



```
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0064 - acc: 0.9981 -
val_loss: 0.0852 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0056 - acc: 0.9982 -
val_loss: 0.1029 - val_acc: 0.9805
```

In [0]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.10294274219236926
Test accuracy: 0.9805
```



In [0]:

```
w_after = model_relu.get_weights()
```

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Batch-Norm on hidden Layers + AdamOptimizer

In [0]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

```

```

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [0]:

```

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 138us/step - loss: 0.3036 - acc: 0.9104 - val_loss: 0.2116 - val_acc: 0.9376
Epoch 2/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1747 - acc: 0.9483 - val_loss: 0.1670 - val_acc: 0.9505
Epoch 3/20
60000/60000 [=====] - 13s 220us/step - loss: 0.1367 - acc: 0.9599 - val_loss: 0.1451 - val_acc: 0.9567
Epoch 4/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1134 - acc: 0.9666 - val_loss: 0.1335 - val_acc: 0.9603
Epoch 5/20
60000/60000 [=====] - 13s 211us/step - loss: 0.0949 - acc: 0.9703 - val_loss: 0.1325 - val_acc: 0.9589
Epoch 6/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0802 - acc: 0.9758 - val_loss: 0.1139 - val_acc: 0.9652
Epoch 7/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0682 - acc: 0.9787 - val_loss: 0.1136 - val_acc: 0.9666
Epoch 8/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0608 - acc: 0.9815 - val_loss: 0.1114 - val_acc: 0.9666
Epoch 9/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0532 - acc: 0.9837 - val_loss: 0.1167 - val_acc: 0.9666
Epoch 10/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0455 - acc: 0.9856 - val_loss: 0.0962 - val_acc: 0.9718
Epoch 11/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0376 - acc: 0.9880 - val_loss: 0.1102 - val_acc: 0.9673
Epoch 12/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0350 - acc: 0.9889 - val_loss: 0.1033 - val_acc: 0.9710
Epoch 13/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0308 - acc: 0.9903 - val_loss: 0.1020 - val_acc: 0.9712

```

```

Epoch 14/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0271 - acc: 0.9913 -
val_loss: 0.1038 - val_acc: 0.9727
Epoch 15/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0231 - acc: 0.9926 -
val_loss: 0.1019 - val_acc: 0.9717
Epoch 16/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0220 - acc: 0.9928 -
val_loss: 0.1110 - val_acc: 0.9703
Epoch 17/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0229 - acc: 0.9928 -
val_loss: 0.1067 - val_acc: 0.9739
Epoch 18/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0203 - acc: 0.9935 -
val_loss: 0.0982 - val_acc: 0.9738
Epoch 19/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0171 - acc: 0.9944 -
val_loss: 0.1056 - val_acc: 0.9706
Epoch 20/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0146 - acc: 0.9952 - val_l
oss: 0.1046 - val_acc: 0.9732

```

In [0]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10456635547156475

Test accuracy: 0.9732





In [0]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [0]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [0]:

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 227us/step - loss: 0.6612 - acc: 0.7951 - val_loss: 0.2860 - val_acc: 0.9166
Epoch 2/20
60000/60000 [=====] - 8s 136us/step - loss: 0.4250 - acc: 0.8710 - val_loss: 0.2545 - val_acc: 0.9252
Epoch 3/20
60000/60000 [=====] - 12s 198us/step - loss: 0.3841 - acc: 0.8846 - val_loss: 0.2391 - val_acc: 0.9298
Epoch 4/20
60000/60000 [=====] - 8s 138us/step - loss: 0.3551 - acc: 0.8927 - val_loss: 0.2279 - val_acc: 0.9325
Epoch 5/20
60000/60000 [=====] - 7s 123us/step - loss: 0.3355 - acc: 0.8986 - val_loss: 0.2127 - val_acc: 0.9356
Epoch 6/20
60000/60000 [=====] - 8s 136us/step - loss: 0.3234 - acc: 0.9031 - val_loss: 0.2029 - val_acc: 0.9387: 1s - loss:
Epoch 7/20
60000/60000 [=====] - 8s 131us/step - loss: 0.3068 - acc: 0.9077 - val_loss: 0.1927 - val_acc: 0.9421
Epoch 8/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2933 - acc: 0.9113 - val_loss: 0.1836 - val_acc: 0.9453
Epoch 9/20
60000/60000 [=====] - 13s 222us/step - loss: 0.2850 - acc: 0.9131 - val_loss: 0.1750 - val_acc: 0.9479

```

```

00000/00000 [-----] 10s 222us/step - loss: 0.2000 - acc: 0.9191 - val_loss: 0.1797 - val_acc: 0.9451
Epoch 10/20
60000/60000 [=====] - 14s 236us/step - loss: 0.2715 - acc: 0.9187 - val_loss: 0.1738 - val_acc: 0.9465
Epoch 11/20
60000/60000 [=====] - 8s 141us/step - loss: 0.2611 - acc: 0.9214 - val_loss: 0.1671 - val_acc: 0.9506
Epoch 12/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2464 - acc: 0.9252 - val_loss: 0.1554 - val_acc: 0.9525
Epoch 13/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2382 - acc: 0.9278 - val_loss: 0.1479 - val_acc: 0.9554
Epoch 14/20
60000/60000 [=====] - 8s 136us/step - loss: 0.2275 - acc: 0.9313 - val_loss: 0.1375 - val_acc: 0.9580
Epoch 15/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2183 - acc: 0.9337 - val_loss: 0.1326 - val_acc: 0.9599
Epoch 16/20
60000/60000 [=====] - 8s 138us/step - loss: 0.2068 - acc: 0.9384 - val_loss: 0.1297 - val_acc: 0.9613 loss: 0.2066 - acc: 0.9384
Epoch 17/20
60000/60000 [=====] - 8s 139us/step - loss: 0.2011 - acc: 0.9395 - val_loss: 0.1181 - val_acc: 0.9646
Epoch 18/20
60000/60000 [=====] - 8s 137us/step - loss: 0.1886 - acc: 0.9435 - val_loss: 0.1145 - val_acc: 0.9658
Epoch 19/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1821 - acc: 0.9451 - val_loss: 0.1104 - val_acc: 0.9662
Epoch 20/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1739 - acc: 0.9473 - val_loss: 0.1093 - val_acc: 0.9679

```

In [0]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.1093290721397847

Test accuracy: 0.9679





In [0]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hidden Layer 1

Hidden Layer 2

Output Layer

Hyper-parameter tuning of Keras models using Sklearn

Assignment

1. Architecture change with 2 hidden layers
2. Architecture change with 3 hidden layers
3. Architecture change with 5 hidden layers

Task 1

In [0]:

```
from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparametersforTask1(activ):

    model = Sequential()
    model.add(Dense(456, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(102, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

Task 2

In [0]:

```
from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparametersforTask2(activ):

    model = Sequential()
    model.add(Dense(484, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(324, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

Task 3

In [0]:

```

from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparametersforTask3(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(456, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(348, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(256, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(124, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model

```

Example code

In [0]:

```

# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)

```

In [0]:

```

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}

```

Task 1 execution

In [0]:

```
in [33]:
```

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.layers.normalization import BatchNormalization
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.layers import Dropout

model = KerasClassifier(build_fn=best_hyperparametersforTask1, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.
```

```
warnings.warn(CV_WARNING, FutureWarning)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
```

```
Instructions for updating:
```

```
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
```

```
Instructions for updating:
```

```
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.global_variables_initializer is deprecated. Please use tf.compat.v1.global_variables_initializer instead.
```

packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.

In [34]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.978167 using {'activ': 'relu'}
0.970483 (0.000768) with: {'activ': 'sigmoid'}
0.978167 (0.001281) with: {'activ': 'relu'}
```

In [36]:

```
from keras.layers.normalization import BatchNormalization

model_task1 = Sequential()

model_task1.add(Dense(456, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_task1.add(BatchNormalization())

model_task1.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_task1.add(BatchNormalization())

model_task1.add(Dense(output_dim, activation='softmax'))

model_task1.summary()
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 456)	357960
batch_normalization_17 (Batch Normalization)	(None, 456)	1824
dense_27 (Dense)	(None, 102)	46614
batch_normalization_18 (Batch Normalization)	(None, 102)	408
dense_28 (Dense)	(None, 10)	1030
Total params: 407,836		
Trainable params: 406,720		
Non-trainable params: 1,116		

In [37]:

```
model_task1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_task1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1970 - acc: 0.9424 - val_loss: 0.1066 - val_acc: 0.9681
Epoch 2/20
```

```

60000/60000 [=====] - 5s 86us/step - loss: 0.0753 - acc: 0.9778 -
val_loss: 0.0932 - val_acc: 0.9711
Epoch 3/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0477 - acc: 0.9859 -
val_loss: 0.0774 - val_acc: 0.9763
Epoch 4/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0359 - acc: 0.9890 -
val_loss: 0.0763 - val_acc: 0.9763
Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0256 - acc: 0.9923 -
val_loss: 0.0717 - val_acc: 0.9779
Epoch 6/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0220 - acc: 0.9931 -
val_loss: 0.0808 - val_acc: 0.9747
Epoch 7/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0188 - acc: 0.9941 -
val_loss: 0.0764 - val_acc: 0.9790
Epoch 8/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0148 - acc: 0.9951 -
val_loss: 0.0696 - val_acc: 0.9802
Epoch 9/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0120 - acc: 0.9965 -
val_loss: 0.0767 - val_acc: 0.9784
Epoch 10/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0146 - acc: 0.9950 -
val_loss: 0.0753 - val_acc: 0.9799
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0123 - acc: 0.9961 -
val_loss: 0.0770 - val_acc: 0.9792
Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0105 - acc: 0.9968 -
val_loss: 0.0730 - val_acc: 0.9791
Epoch 13/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0096 - acc: 0.9968 -
val_loss: 0.0763 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0093 - acc: 0.9970 -
val_loss: 0.0893 - val_acc: 0.9780
Epoch 15/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0093 - acc: 0.9968 -
val_loss: 0.0776 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0083 - acc: 0.9974 -
val_loss: 0.0764 - val_acc: 0.9814
Epoch 17/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0060 - acc: 0.9982 -
val_loss: 0.0909 - val_acc: 0.9795
Epoch 18/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0073 - acc: 0.9975 -
val_loss: 0.0904 - val_acc: 0.9800
Epoch 19/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0077 - acc: 0.9973 -
val_loss: 0.0759 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0060 - acc: 0.9980 -
val_loss: 0.0806 - val_acc: 0.9804

```

In [77]:

```

import mpld3
from mpld3 import plugins

score_task1 = model_task1.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score_task1[0])
print('Test accuracy:', score_task1[1])

ax_task1.set_xlabel('epoch') ; ax_task1.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x_task1 = list(range(1,nb_epoch+1))

vy_task1 = history.history['val_loss']
ty_task1 = history.history['loss']

print(vy_task1)

```

```
print(ty_task1)

plt_dynamic_plot(x_task1, vy_task1, ty_task1, ax_task1)
mpld3.display()
```

```
Test score: 0.08062241236004375
Test accuracy: 0.9804
[0.10663776563256978, 0.09317723634988069, 0.07740170887410641, 0.07626734635229222,
0.07168362509161234, 0.08082462112847716, 0.07638109823968262, 0.06958439976067748,
0.07670358981615864, 0.07533570332399103, 0.07699054275283124, 0.07295610949172406,
0.07633888032212853, 0.08933395910579711, 0.07756036461495677, 0.07637776467328658,
0.09086646485980746, 0.09044575633805507, 0.07594511643507867, 0.08062241120594554]
[0.19695996968746185, 0.07529902205864589, 0.0477438554495573, 0.035873308963576954,
0.025566030979156495, 0.022009523358444374, 0.018790008194496235, 0.014786293017243346,
0.012021565819283327, 0.014606639272719622, 0.012285229281087717, 0.010518471387246002,
0.009619418952443327, 0.009341397868438314, 0.009349660067384441, 0.008259907729923726,
0.006044347383423398, 0.007345158197854956, 0.0076754562556743625, 0.006027616090932861]
```

Out[77]:

Task 2 execution

In [79]:

```
activ = ['sigmoid','relu']

from keras.layers.normalization import BatchNormalization
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.layers import Dropout

model2 = KerasClassifier(build_fn=best_hyperparametersforTask2, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid_task2 = GridSearchCV(estimator=model2, param_grid=param_grid)
grid_result2 = grid_task2.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The
default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this
warning.
    warnings.warn(CV_WARNING, FutureWarning)
```

In [80]:

```
print("Best: %f using %s" % (grid_result2.best_score_, grid_result2.best_params_))
means = grid_result2.cv_results_['mean_test_score']
stds = grid_result2.cv_results_['std_test_score']
params = grid_result2.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.977517 using {'activ': 'relu'}
0.970433 (0.001496) with: {'activ': 'sigmoid'}
0.977517 (0.000931) with: {'activ': 'relu'}
```

In [82]:

```
from keras.layers.normalization import BatchNormalization

model_task2 = Sequential()

model_task2.add(Dense(484, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_task2.add(BatchNormalization())
```

```

model_task2.add(Dropout(0.5))
model_task2.add(Dense(324, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task2.add(BatchNormalization())
model_task2.add(Dropout(0.5))
model_task2.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task2.add(BatchNormalization())
model_task2.add(Dropout(0.5))
model_task2.add(Dense(output_dim, activation='softmax'))

model_task2.summary()

```

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_60 (Dense)	(None, 484)	379940
batch_normalization_42 (Batch Normalization)	(None, 484)	1936
dropout_38 (Dropout)	(None, 484)	0
dense_61 (Dense)	(None, 324)	157140
batch_normalization_43 (Batch Normalization)	(None, 324)	1296
dropout_39 (Dropout)	(None, 324)	0
dense_62 (Dense)	(None, 128)	41600
batch_normalization_44 (Batch Normalization)	(None, 128)	512
dropout_40 (Dropout)	(None, 128)	0
dense_63 (Dense)	(None, 10)	1290
Total params: 583,714		
Trainable params: 581,842		
Non-trainable params: 1,872		

In [83]:

```

model_task2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history2 = model_task2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 12s 202us/step - loss: 0.6564 - acc: 0.7996 - val_loss: 0.1881 - val_acc: 0.9407
Epoch 2/20
60000/60000 [=====] - 7s 118us/step - loss: 0.2767 - acc: 0.9178 - val_loss: 0.1322 - val_acc: 0.9582
Epoch 3/20
60000/60000 [=====] - 7s 117us/step - loss: 0.2130 - acc: 0.9364 - val_loss: 0.1096 - val_acc: 0.9663
Epoch 4/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1776 - acc: 0.9472 - val_loss: 0.1004 - val_acc: 0.9694
Epoch 5/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1581 - acc: 0.9528 - val_loss: 0.0898 - val_acc: 0.9735
Epoch 6/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1379 - acc: 0.9582 - val_loss: 0.0820 - val_acc: 0.9749
Epoch 7/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1308 - acc: 0.9608 - val_loss: 0.0794 - val_acc: 0.9760
Epoch 8/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1180 - acc: 0.9654 - val_loss: 0.0776 - val_acc: 0.9758
Epoch 9/20

```

```

Epoch 9/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1101 - acc: 0.9668 -
val_loss: 0.0803 - val_acc: 0.9776
Epoch 10/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1000 - acc: 0.9702 -
val_loss: 0.0722 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0989 - acc: 0.9698 -
val_loss: 0.0686 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0867 - acc: 0.9732 -
val_loss: 0.0682 - val_acc: 0.9806
Epoch 13/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0855 - acc: 0.9740 -
val_loss: 0.0670 - val_acc: 0.9804
Epoch 14/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0774 - acc: 0.9763 -
val_loss: 0.0649 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0781 - acc: 0.9760 -
val_loss: 0.0643 - val_acc: 0.9814
Epoch 16/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0748 - acc: 0.9772 -
val_loss: 0.0682 - val_acc: 0.9818
Epoch 17/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0711 - acc: 0.9781 -
val_loss: 0.0662 - val_acc: 0.9799
Epoch 18/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0720 - acc: 0.9773 -
val_loss: 0.0627 - val_acc: 0.9831
Epoch 19/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0668 - acc: 0.9789 -
val_loss: 0.0583 - val_acc: 0.9827
Epoch 20/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0632 - acc: 0.9801 -
val_loss: 0.0645 - val_acc: 0.9825

```

In [84]:

```

score_task2 = model_task2.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score_task2[0])
print('Test accuracy:', score_task2[1])

ax_task1.set_xlabel('epoch') ; ax_task1.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x_task1 = list(range(1,nb_epoch+1))

vy_task1 = history2.history['val_loss']
ty_task1 = history2.history['loss']

print(vy_task1)
print(ty_task1)

plt_dynamic_plot(x_task1, vy_task1, ty_task1, ax_task1)
mpld3.display()

```

```

Test score: 0.06452691478281632
Test accuracy: 0.9825
[0.18811061244010926, 0.1321807575829327, 0.1096077186241746, 0.10037248664498329,
0.08980758393295109, 0.08201991958636791, 0.0793870957940817, 0.07755979443769902,
0.08031285183485598, 0.07217185893226415, 0.06863237277567387, 0.0682252861815039,
0.06699350098753348, 0.06486856495265383, 0.06426630284576677, 0.0682348092280794,
0.06615413975138217, 0.06274196430703159, 0.058282991545554254, 0.06452691504568793]
[0.6564425819079082, 0.27673976113796234, 0.21299197249412535, 0.17755161137183506,
0.1581129201332728, 0.1378718542456627, 0.13079650530020395, 0.11802167086998622,
0.11007206436594327, 0.09997802500327428, 0.09885259782274564, 0.08671309248606364,
0.08551792087952297, 0.07736172651052475, 0.07810872188607852, 0.07475954596996308,
0.07110208432475726, 0.07195941407283148, 0.06680011056760947, 0.063217876030008]

```

Out[84]:

Task 3 Execution

In [86]:

```
activ = ['sigmoid', 'relu']

from keras.layers.normalization import BatchNormalization
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.layers import Dropout

model3 = KerasClassifier(build_fn=best_hyperparametersforTask3, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid_task3 = GridSearchCV(estimator=model3, param_grid=param_grid)
grid_result3 = grid_task3.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The
default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this
warning.
    warnings.warn(CV_WARNING, FutureWarning)
```

In [87]:

```
print("Best: %f using %s" % (grid_result3.best_score_, grid_result3.best_params_))
means = grid_result3.cv_results_['mean_test_score']
stds = grid_result3.cv_results_['std_test_score']
params = grid_result3.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.976133 using {'activ': 'relu'}
0.965667 (0.001093) with: {'activ': 'sigmoid'}
0.976133 (0.001047) with: {'activ': 'relu'}
```

In [90]:

```
model_task3 = Sequential()

model_task3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_task3.add(BatchNormalization())
model_task3.add(Dropout(0.5))
model_task3.add(Dense(456, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task3.add(BatchNormalization())
model_task3.add(Dropout(0.5))
model_task3.add(Dense(348, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task3.add(BatchNormalization())
model_task3.add(Dropout(0.5))
model_task3.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task3.add(BatchNormalization())
model_task3.add(Dropout(0.5))
model_task3.add(Dense(124, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_task3.add(BatchNormalization())
model_task3.add(Dropout(0.5))
model_task3.add(Dense(output_dim, activation='softmax'))

model_task3.summary()
```

Model: "sequential_38"

Layer (type)	Output Shape	Param #
=====		

dense_139 (Dense)	(None, 512)	401920
batch_normalization_106 (Batch Normalization)	(None, 512)	2048
dropout_102 (Dropout)	(None, 512)	0
dense_140 (Dense)	(None, 456)	233928
batch_normalization_107 (Batch Normalization)	(None, 456)	1824
dropout_103 (Dropout)	(None, 456)	0
dense_141 (Dense)	(None, 348)	159036
batch_normalization_108 (Batch Normalization)	(None, 348)	1392
dropout_104 (Dropout)	(None, 348)	0
dense_142 (Dense)	(None, 256)	89344
batch_normalization_109 (Batch Normalization)	(None, 256)	1024
dropout_105 (Dropout)	(None, 256)	0
dense_143 (Dense)	(None, 124)	31868
batch_normalization_110 (Batch Normalization)	(None, 124)	496
dropout_106 (Dropout)	(None, 124)	0
dense_144 (Dense)	(None, 10)	1250
=====		
Total params: 924,130		
Trainable params: 920,738		
Non-trainable params: 3,392		

In [91]:

```
model_task3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history3 = model_task3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 23s 390us/step - loss: 1.1447 - acc: 0.6384 - val_loss: 0.2758 - val_acc: 0.9174
Epoch 2/20
60000/60000 [=====] - 10s 170us/step - loss: 0.4032 - acc: 0.8816 - val_loss: 0.1789 - val_acc: 0.9472
Epoch 3/20
60000/60000 [=====] - 10s 168us/step - loss: 0.2911 - acc: 0.9179 - val_loss: 0.1489 - val_acc: 0.9575
Epoch 4/20
60000/60000 [=====] - 10s 171us/step - loss: 0.2391 - acc: 0.9324 - val_loss: 0.1174 - val_acc: 0.9652
Epoch 5/20
60000/60000 [=====] - 10s 168us/step - loss: 0.2065 - acc: 0.9418 - val_loss: 0.1093 - val_acc: 0.9695
Epoch 6/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1834 - acc: 0.9487 - val_loss: 0.1016 - val_acc: 0.9715
Epoch 7/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1627 - acc: 0.9542 - val_loss: 0.0974 - val_acc: 0.9731
Epoch 8/20
60000/60000 [=====] - 10s 168us/step - loss: 0.1483 - acc: 0.9581 - val_loss: 0.0838 - val_acc: 0.9758
Epoch 9/20
60000/60000 [=====] - 10s 168us/step - loss: 0.1375 - acc: 0.9604 - val_loss: 0.0855 - val_acc: 0.9757
Epoch 10/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1313 - acc: 0.9628 - val_loss: 0.0840 - val_acc: 0.9772
```

```

Epoch 11/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1233 - acc: 0.9659 - val_loss: 0.0799 - val_acc: 0.9780
Epoch 12/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1184 - acc: 0.9660 - val_loss: 0.0814 - val_acc: 0.9772
Epoch 13/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1148 - acc: 0.9685 - val_loss: 0.0722 - val_acc: 0.9809
Epoch 14/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1076 - acc: 0.9695 - val_loss: 0.0735 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1047 - acc: 0.9703 - val_loss: 0.0677 - val_acc: 0.9808
Epoch 16/20
60000/60000 [=====] - 10s 172us/step - loss: 0.1014 - acc: 0.9719 - val_loss: 0.0690 - val_acc: 0.9817
Epoch 17/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0950 - acc: 0.9736 - val_loss: 0.0694 - val_acc: 0.9818
Epoch 18/20
60000/60000 [=====] - 10s 170us/step - loss: 0.0910 - acc: 0.9749 - val_loss: 0.0644 - val_acc: 0.9836
Epoch 19/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0930 - acc: 0.9741 - val_loss: 0.0714 - val_acc: 0.9819
Epoch 20/20
60000/60000 [=====] - 10s 169us/step - loss: 0.0826 - acc: 0.9762 - val_loss: 0.0692 - val_acc: 0.9825

```

In [92]:

```

score_task3 = model_task3.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score_task3[0])
print('Test accuracy:', score_task3[1])

ax_task1.set_xlabel('epoch') ; ax_task1.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x_task1 = list(range(1,nb_epoch+1))

vy_task1 = history3.history['val_loss']
ty_task1 = history3.history['loss']

print(vy_task1)
print(ty_task1)

plt_dynamic_plot(x_task1, vy_task1, ty_task1, ax_task1)
mpld3.display()

```

```

Test score: 0.06919898579865694
Test accuracy: 0.9825
[0.2757951403617859, 0.17889915275275708, 0.14887760931998492, 0.11736068408563734,
0.10926629482582212, 0.10164899545609951, 0.09743993790000677, 0.08383604726046323,
0.08553117757583968, 0.08398494211179204, 0.07987379637137056, 0.08139870572201907,
0.07221250504720957, 0.07348079082481564, 0.06771763956397772, 0.06896937229735776,
0.06935005094539375, 0.06440088580437005, 0.07141120507828891, 0.06919898539423011]
[1.144697232834498, 0.40316257503827413, 0.29111041218439737, 0.23914790218671164,
0.20650551250775656, 0.18335089112122854, 0.16265787526369094, 0.14834913578828177,
0.13748974556128185, 0.13125558004975318, 0.12330681715607643, 0.11839671297967434,
0.11480624519586563, 0.10761370351711909, 0.10465052956342698, 0.10139177598158519,
0.09499700264135996, 0.09102033909161886, 0.09296152342955272, 0.08258540489971637]

```

Out [92]:

Conclusion

In [93]:

```

from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Architecture", "Activation", "Optimizer", "Test score", "Test Accuracy"]

x.add_row(["2 Layer NN", "ReLU", "Adam", 0.0806, 0.9804 ])
x.add_row(["3 Layer NN", "ReLU", "Adam", 0.0645, 0.9825])
x.add_row(["5 Layer NN", "ReLU", "Adam", 0.0691, 0.9825])

print(x)

```

```

+-----+-----+-----+-----+-----+
| Architecture | Activation | Optimizer | Test score | Test Accuracy |
+-----+-----+-----+-----+-----+
| 2 Layer NN | ReLU | Adam | 0.0806 | 0.9804 |
| 3 Layer NN | ReLU | Adam | 0.0645 | 0.9825 |
| 5 Layer NN | ReLU | Adam | 0.0691 | 0.9825 |
+-----+-----+-----+-----+-----+

```